

# Project 2: UNIX Shell Programming and Linux Kernel Module for Task Information

CS307-Operating System (D), CS356-Operating System Course Design, Chentao Wu, Spring 2020.

Name: 方泓杰(Hongjie Fang) Student ID: 518030910150 Email: galaxies@sjtu.edu.cn

## 1 UNIX Shell Programming

According to the project descriptions in textbook, we are required to write a UNIX Shell which provides the following functions.

- Creating the child process and executing the command in the child;
- Providing a history feature;
- Adding support of input and output redirection;
- Allowing the parent and child processes to communicate via a pipe.

**Solution.** Here is my method to implement the simple UNIX Shell program (`shell.c`).

- First, we implement a `standardlize_inst` function to standardlize the instruction in a standard form, which means there is no extra space and tab in the instruction.
- After standardlization, we can check whether this instruction asks for concurrent execution. If true, then we set the `concurrent` flag to 1, and the parent process do not have to execute `wait(NULL)` instruction to wait for its child process to finish execution since we can execute two processes concurrently.
- Then we can handle some special instructions such as `exit` and `!!`.
  - For `exit` instruction, we simply set the `should_run` flag to false, then it will jump out of the loop and finish executing the main function;
  - For `!!` instruction, we store the last instruction in the variable `last_inst` so we only need to replace the current instruction `inst` with the last instruction `last_inst`. To handle the error condition, we need to set a `have_last_inst` flag to check whether the current instruction do not have the previous instruction in the history.
- After that, we create a child process using `fork()` function, and implement a `parse` function in child process to parse the instruction and return several arguments `args[]` of the instruction. In parent process, we actually do nothing except waiting (if necessary).
- After parsing instruction, we can check whether the instruction needs a pipe. If true, then we can use `pipe(pipe_fd)` function to generate a pipe and use it to implement data transfer. The implementation of pipe transfer is similar to Section 3.7.4 in textbook.
- If the instruction does not need a pipe, then we can find out the input/output redirections of the instruction by examining the result of parsing. If the instruction needs redirections, then we can redirect the input/output using `dup2(fd, ...)` function.
- Then, we can execute the instruction using the `execvp(...)` function and the arguments `args[]` we have parsed before.

- After execution, we need to release the spaces of variables, close the files, and then we can exit from child process.

Here is the specific implementation of the UNIX Shell.

```
1 # include <stdio.h>
2 # include <fcntl.h>
3 # include <stdlib.h>
4 # include <string.h>
5 # include <unistd.h>
6 # include <sys/wait.h>
7 # include <sys/types.h>
8
9 # define MAX_LINE 80
10 # define READ_END 0
11 # define WRITE_END 1
12
13 // Clear the string.
14 void clear_str(char *str);
15
16 // Check whether the instruction is concurrent.
17 int check_concurrent(char *inst);
18
19 // Standardize the instruction.
20 void standardize_inst(char *inst);
21
22 // Parse the instruction
23 int parse(char *inst, char **args);
24
25 // Debug program for parsing.
26 void debug_parse(char *args[], int argn);
27
28 int main(void) {
29     // arguments, instruction, last instruction
30     char *args[MAX_LINE / 2 + 1], *inst, *last_inst;
31     // whether have the last instruction, cocurrent status
32     int have_last_inst = 0, concurrent = 0;
33     // input redirect filename, output redirect filename
34     char *in_file, *out_file;
35
36     inst = (char*) malloc(MAX_LINE * sizeof(char));
37     last_inst = (char*) malloc(MAX_LINE * sizeof(char));
38     in_file = (char*) malloc(MAX_LINE * sizeof(char));
39     out_file = (char*) malloc(MAX_LINE * sizeof(char));
40
41     clear_str(last_inst);
42
43     int should_run = 1;
44
45     pid_t pid;
46
```

```

47 while(should_run) {
48     printf("osh> ");
49     fflush(stdout);
50     if (concurrent) wait(NULL);
51
52     concurrent = 0;
53     clear_str(inst);
54
55     fgets(inst, MAX_LINE, stdin);
56
57     standardlize_inst(inst);
58     concurrent = check_concurrent(inst);
59
60     // exit shell
61     if (strcmp(inst, "exit") == 0) {
62         should_run = 0;
63         continue;
64     }
65
66     // execute the last instruction
67     if (strcmp(inst, "!!") == 0) {
68         if (have_last_inst == 0) {
69             fprintf(stderr, "Error: No commands in history.\n");
70             continue;
71         } else {
72             printf("%s\n", last_inst);
73             strcpy(inst, last_inst);
74         }
75     }
76
77     pid = fork();
78     if (pid < 0) fprintf(stderr, "Error: Fork failed!\n");
79     else {
80         if (pid == 0) {
81             // child process
82             // whether an error has occurred
83             int error_occur = 0;
84
85             // allocate space for commands & arguments
86             for (int i = 0; i <= MAX_LINE / 2; ++ i)
87                 args[i] = (char*) malloc(MAX_LINE * sizeof(char));
88
89             // parse the instruction
90             int argn = parse(inst, args);
91
92             // free the space of extra commands & extra arguments
93             for (int i = argn; i <= MAX_LINE / 2; ++ i) {
94                 free(args[i]);
95                 args[i] = NULL;
96             }

```

```

97     if (concurrent == 1) {
98         -- argn;
99         free(args[argn]);
100        args[argn] = NULL;
101    }
102
103    // find pipe
104    int pipe_pos = -1;
105    for (int i = 0; i < argn; ++ i)
106        if (strcmp(args[i], "|") == 0) {
107            pipe_pos = i;
108            break;
109        }
110
111    if(pipe_pos >= 0) {
112        // pipe found
113        if (pipe_pos == 0 || pipe_pos == argn - 1) {
114            fprintf(stderr, "Error: Unexpected syntax '|'.\n");
115            error_occur = 1;
116        }
117
118        // pipe fd
119        int pipe_fd[2];
120
121        if (pipe(pipe_fd) == -1) {
122            fprintf(stderr, "Error: Pipe Failed!\n");
123            error_occur = 1;
124        }
125
126        if(error_occur == 0) {
127            // fork a grandson process
128            pid = fork();
129            if (pid < 0) {
130                fprintf(stderr, "Error: Fork failed!\n");
131                error_occur = 1;
132            } else {
133                if (pid == 0) {
134                    // grandchild process
135                    for (int i = pipe_pos; i < argn; ++ i) {
136                        free(args[i]);
137                        args[i] = NULL;
138                    }
139                    argn = pipe_pos;
140
141                    close(pipe_fd[READ_END]);
142                    if (error_occur == 0 && dup2(pipe_fd[WRITE_END], STDOUT_FILENO) <
143                        0) {
144                        fprintf(stderr, "Error: dup2 failed!\n");
145                        error_occur = 1;
146                    }

```

```

146
147         if(error_occur == 0 && argn > 0) execvp(args[0], args);
148
149         close(pipe_fd[WRITE_END]);
150
151         // free the spaces
152         for (int i = 0; i < argn; ++ i) free(args[i]);
153         free(inst);
154         free(last_inst);
155         free(in_file);
156         free(out_file);
157
158         exit(error_occur);
159     } else {
160         // child process
161         wait(NULL);
162         for (int i = 0; i <= pipe_pos; ++ i) free(args[i]);
163         for (int i = pipe_pos + 1; i < argn; ++ i) args[i - pipe_pos - 1] =
            args[i];
164         for (int i = argn - pipe_pos - 1; i < argn; ++ i) args[i] = NULL;
165         argn = argn - pipe_pos - 1;
166
167         close(pipe_fd[WRITE_END]);
168         if (error_occur == 0 && dup2(pipe_fd[READ_END], STDIN_FILENO) < 0)
169         {
170             fprintf(stderr, "Error: dup2 failed!\n");
171             error_occur = 1;
172         }
173
174         if(error_occur == 0 && argn > 0) execvp(args[0], args);
175
176         close(pipe_fd[READ_END]);
177     }
178 }
179 } else {
180     // find in_redirect or out_redirect
181     int in_redirect = 0, out_redirect = 0, in_fd = -1, out_fd = -1;
182     while (argn >= 2 && (strcmp(args[argn - 2], "<") == 0 || strcmp(args[
        argn - 2], ">") == 0)) {
183         argn -= 2;
184         if (strcmp(args[argn], "<") == 0) {
185             in_redirect = 1;
186             strcpy(in_file, args[argn + 1]);
187         } else {
188             out_redirect = 1;
189             strcpy(out_file, args[argn + 1]);
190         }
191         free(args[argn]); args[argn] = NULL;
192         free(args[argn + 1]); args[argn + 1] = NULL;

```

```

193     }
194
195     // redirect input
196     if (error_occur == 0 && in_redirect == 1) {
197         in_fd = open(in_file, O_RDONLY, 0644);
198         if (error_occur == 0 && in_fd < 0) {
199             fprintf(stderr, "Error: No such files.\n");
200             error_occur = 1;
201         }
202         if (error_occur == 0 && dup2(in_fd, STDIN_FILENO) < 0) {
203             fprintf(stderr, "Error: dup2 failed!\n");
204             error_occur = 1;
205         }
206     }
207
208     // redirect output
209     if (error_occur == 0 && out_redirect == 1) {
210         out_fd = open(out_file, O_WRONLY | O_TRUNC | O_CREAT, 0644);
211         if (error_occur == 0 && out_fd < 0) {
212             fprintf(stderr, "Error: No such files.\n");
213             error_occur = 1;
214         }
215         if (error_occur == 0 && dup2(out_fd, STDOUT_FILENO) < 0) {
216             fprintf(stderr, "Error: dup2 failed!\n");
217             error_occur = 1;
218         }
219     }
220
221     // not an empty instruction & no error occur, then execute the instruction
222     if (error_occur == 0 && argn != 0)
223         execvp(args[0], args);
224
225     // close the files
226     if (in_redirect == 1 && in_fd > 0) close(in_fd);
227     if (out_redirect == 1 && out_fd > 0) close(out_fd);
228 }
229
230 // free the spaces
231 for (int i = 0; i < argn; ++ i) free(args[i]);
232 free(inst);
233 free(last_inst);
234 free(in_file);
235 free(out_file);
236
237 // child process exit
238 exit(error_occur);
239 } else {
240     // parent process
241     if(concurrent == 0) wait(NULL);
242 }

```

```

243     }
244
245     if(have_last_inst == 0) have_last_inst = 1;
246     strcpy(last_inst, inst);
247 }
248
249 // free the spaces
250 free(inst);
251 free(last_inst);
252 free(in_file);
253 free(out_file);
254 return 0;
255 }
256
257
258 // Clear the string.
259 void clear_str(char *str) {
260     memset(str, 0, sizeof(str));
261 }
262
263 // Check whether the instruction is concurrent.
264 int check_concurrent(char *inst) {
265     int len = strlen(inst);
266     if(len && inst[len - 1] == '&') return 1;
267     return 0;
268 }
269
270 // Standardize the instruction.
271 // Specific function: clear the extra space & tab & enter in the instruction.
272 void standardize_inst(char *inst) {
273     int len = strlen(inst);
274
275     char *temp = (char*) malloc(len * sizeof(char));
276     for (int i = 0; i < len; ++ i) temp[i] = inst[i];
277     clear_str(inst);
278
279     int new_len = 0, last_blank = 1;
280     for (int i = 0; i < len; ++ i) {
281         if (temp[i] == ' ' || temp[i] == '\n' || temp[i] == '\t') {
282             if (last_blank == 0) {
283                 inst[new_len++] = ' ';
284                 last_blank = 1;
285             }
286         } else {
287             inst[new_len++] = temp[i];
288             last_blank = 0;
289         }
290     }
291     if(inst[new_len - 1] == ' ') inst[new_len - 1] = 0;
292

```

```

293     free(temp);
294 }
295
296 // Parse the instruction .
297 // Specific function: parse the instruction and find out the command & arguments.
298 int parse(char *inst, char **args) {
299     int len = strlen(inst);
300
301     // find out the arguments
302     int argn = 0;
303     for (int i = 0; i < len; ++ i) {
304         clear_str(args[argn]);
305
306         int j = i;
307         while(j < len && inst[j] != ' ') {
308             args[argn][j - i] = inst[j];
309             ++ j;
310         }
311         if ((args[argn][0] == '<' || args[argn][0] == '>' || args[argn][0] == '|') &&
312             j > i+1) {
313             strcpy(args[argn + 1], args[argn] + 1);
314             for (int k = 1; k < j - i; ++ k) args[argn][k] = 0;
315             ++ argn;
316         }
317         i = j;
318         ++ argn;
319     }
320
321     return argn;
322 }
323
324 // Debug program for parsing .
325 void debug_parse(char *args[], int argn) {
326     fprintf(stderr, "Comm: %s, total %d arguments\n", args[0], argn);
327     for (int i = 0; i < argn; ++ i)
328         fprintf(stderr, "args[%d] = %s\n", i, args[i]);

```

We can test the UNIX Shell program by entering the following instructions. These instructions can test every required function of the UNIX Shell.

```

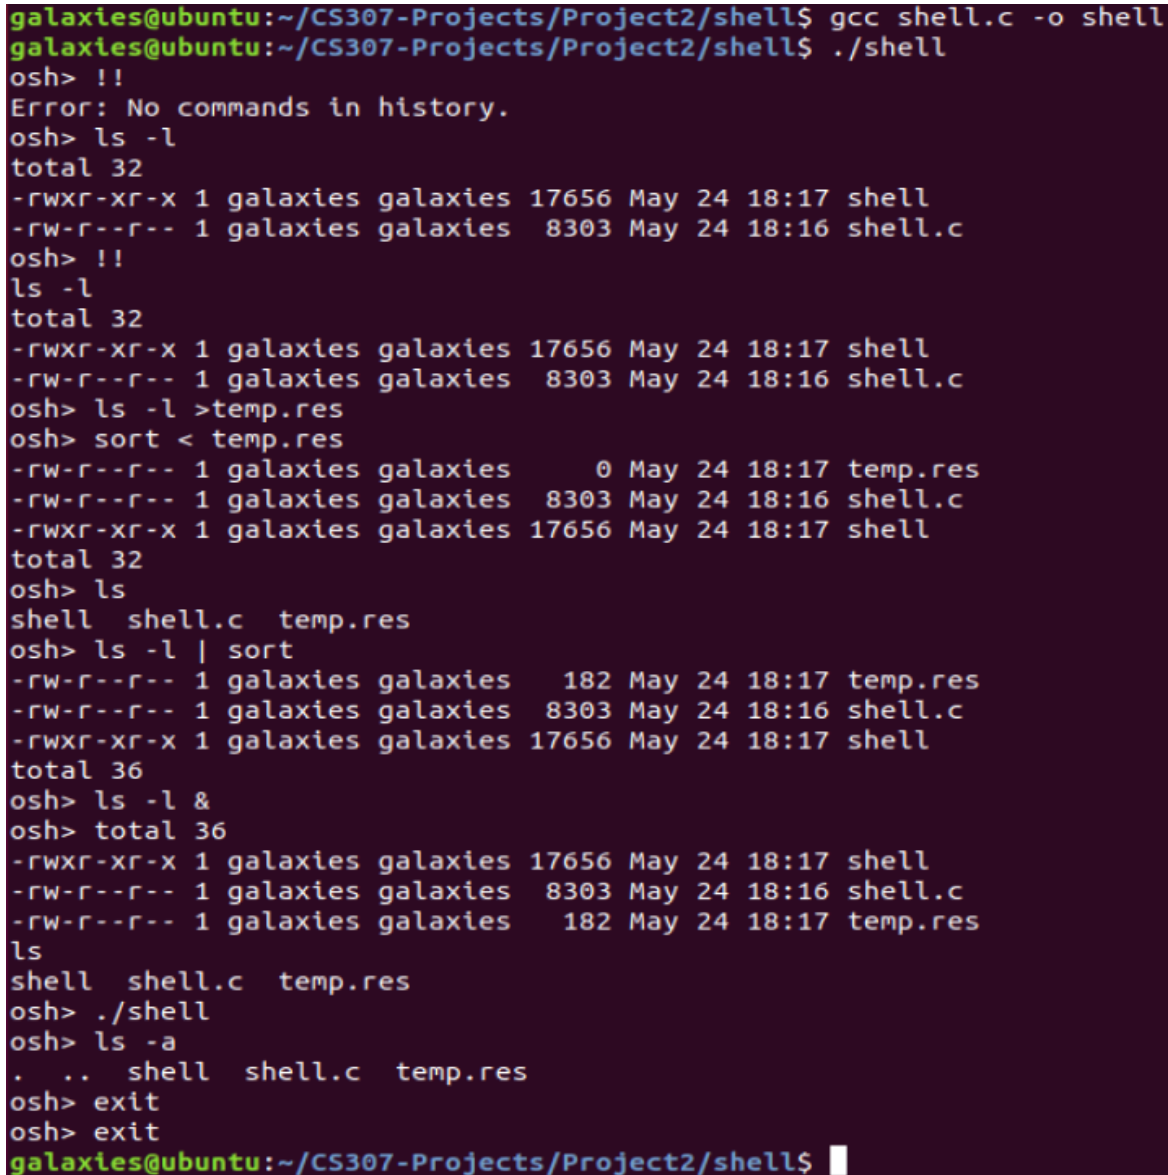
1 gcc shell.c -o ./shell
2 ./shell
3 !!
4 ls -l
5 !!
6 ls -l > temp.res
7 sort < temp.res
8 ls
9 ls -l
10 ls -l | sort

```



```
11 | ls -l &
12 | ls
13 | ./shell
14 | ls -a
15 | exit
16 | exit
```

Here is the execution result of my UNIX Shell program after entering the instructions (Fig. 1).



```
galaxies@ubuntu:~/CS307-Projects/Project2/shell$ gcc shell.c -o shell
galaxies@ubuntu:~/CS307-Projects/Project2/shell$ ./shell
osh> !!
Error: No commands in history.
osh> ls -l
total 32
-rwxr-xr-x 1 galaxies galaxies 17656 May 24 18:17 shell
-rw-r--r-- 1 galaxies galaxies 8303 May 24 18:16 shell.c
osh> !!
ls -l
total 32
-rwxr-xr-x 1 galaxies galaxies 17656 May 24 18:17 shell
-rw-r--r-- 1 galaxies galaxies 8303 May 24 18:16 shell.c
osh> ls -l >temp.res
osh> sort < temp.res
-rw-r--r-- 1 galaxies galaxies 0 May 24 18:17 temp.res
-rw-r--r-- 1 galaxies galaxies 8303 May 24 18:16 shell.c
-rwxr-xr-x 1 galaxies galaxies 17656 May 24 18:17 shell
total 32
osh> ls
shell shell.c temp.res
osh> ls -l | sort
-rw-r--r-- 1 galaxies galaxies 182 May 24 18:17 temp.res
-rw-r--r-- 1 galaxies galaxies 8303 May 24 18:16 shell.c
-rwxr-xr-x 1 galaxies galaxies 17656 May 24 18:17 shell
total 36
osh> ls -l &
osh> total 36
-rwxr-xr-x 1 galaxies galaxies 17656 May 24 18:17 shell
-rw-r--r-- 1 galaxies galaxies 8303 May 24 18:16 shell.c
-rw-r--r-- 1 galaxies galaxies 182 May 24 18:17 temp.res
ls
shell shell.c temp.res
osh> ./shell
osh> ls -a
. .. shell shell.c temp.res
osh> exit
osh> exit
galaxies@ubuntu:~/CS307-Projects/Project2/shell$
```

Figure 1: The execution result of my UNIX Shell program

□

## 2 Linux Kernel Module for Task Information

In this project, we are required to write a Linux kernel module that use the `/proc` file system for displaying a task's information based on its process identifier value `pid`.

**Solution.** Here is my method to implement the required Linux kernel module.

- In `proc_write` function, we can use the `kstrtol()` function to read the value from a string. An important fact is that `usr_buf` may do not have an end-of-string sign `'\0'`, therefore we need to add this sign to the end of the string manually.
- In `proc_write` function, we need to use `kmalloc()` and `kfree()` function to allocate and release memory; they are actually the kernel version of `malloc()` and `free()`.
- In `proc_read` function, we can use the `pid_task()` function to read the information in PCB, and we need to use `find_vpid()` to find the corresponding PCB using its `pid`. What we need are the `comm` and `state` value of the PCB.
- We also perform an error checking for invalid `pid`. This will cause the `pid_task()` function returns `NULL`. We simply report the error and return 0.
- The other parts of the program is very similar to project 1, we can use the similar methods to implement them.

Here is the specific implementation of the Linux kernel module for task information (`pid.c`).

```
1 # include <linux/init.h>
2 # include <linux/slab.h>
3 # include <linux/sched.h>
4 # include <linux/module.h>
5 # include <linux/kernel.h>
6 # include <linux/proc_fs.h>
7 # include <linux/vmalloc.h>
8 # include <linux/uaccess.h>
9 # include <asm/uaccess.h>
10
11 # define BUFFER_SIZE 128
12 # define PROC_NAME "pid"
13
14 // the current pid
15 static long cur_pid;
16
17 static ssize_t proc_read(struct file *file, char *buf, size_t count, loff_t *pos)
18     ;
19
20 static ssize_t proc_write(struct file *file, const char __user *usr_buf, size_t
21     count, loff_t *pos);
22
23 static struct file_operations proc_ops = {
24     .owner = THIS_MODULE,
25     .read = proc_read,
26     .write = proc_write,
27 };
28
29
```

```

26 static int proc_init(void) {
27     /* create /proc files */
28     proc_create(PROC_NAME, 0666, NULL, &proc_ops);
29     printk(KERN_INFO "/proc/" PROC_NAME " is created!\n");
30     return 0;
31 }
32
33 static void proc_exit(void) {
34     /* remove /proc files */
35     remove_proc_entry(PROC_NAME, NULL);
36     printk(KERN_INFO "/proc/" PROC_NAME " is removed!\n");
37 }
38
39 static ssize_t proc_read(struct file *file, char __user *usr_buf, size_t count,
40     loff_t *pos) {
41     int rv = 0;
42     char buffer[BUFFER_SIZE];
43     static int completed = 0;
44     struct task_struct *PCB = NULL;
45
46     if (completed) {
47         completed = 0;
48         return 0;
49     }
50     PCB = pid_task(find_vpid(cur_pid), PIDTYPE_PID);
51     if (PCB == NULL) {
52         printk(KERN_INFO "Invalid PID!\n");
53         return 0;
54     }
55
56     completed = 1;
57
58     rv = sprintf(buffer, "command = [%s] pid = [%ld] state = [%ld]\n", PCB -> comm,
59         cur_pid, PCB -> state);
60     copy_to_user(usr_buf, buffer, rv);
61
62     return rv;
63 }
64
65 static ssize_t proc_write(struct file *file, const char __user *usr_buf, size_t
66     count, loff_t *pos) {
67     char *k_mem;
68
69     // allocate kernel memory
70     k_mem = kmalloc(count, GFP_KERNEL);
71
72     // copies user space usr_buf to kernel buffer
73     if (copy_from_user(k_mem, usr_buf, count)) {

```

```

73     printk(KERN_INFO "Error copying from user\n");
74     return -1;
75 }
76
77 // the end of string k_mem
78 k_mem[count] = 0;
79
80 // extract the number into the variable pid using kstrtoul
81 kstrtoul(k_mem, 10, &cur_pid);
82
83 // free the memory
84 kfree(k_mem);
85
86 return count;
87 }
88
89 module_init(proc_init);
90 module_exit(proc_exit);
91
92 MODULE_LICENSE("GPL");
93 MODULE_DESCRIPTION("Pid Module");
94 MODULE_AUTHOR("Galaxies");

```

And here is the Makefile file of the project.

```

1 obj-m := pid.o
2
3 all:
4     make -C /usr/src/linux-5.5.8/ M=$(shell pwd) modules
5 clean:
6     make -C /usr/src/linux-5.5.8/ M=$(shell pwd) clean

```

We can enter the following instructions to test the kernel module.

```

1 make
2 sudo dmesg -C
3 sudo insmod pid.ko
4 sudo dmesg
5 echo "1395" > /proc/pid
6 cat /proc/pid
7 echo "1" > /proc/pid
8 cat /proc/pid
9 echo "5" > /proc/pid
10 cat /proc/pid
11 echo "6" > /proc/pid
12 cat /proc/pid
13 sudo rmmod pid
14 sudo dmesg

```

Here is the execution result of the program (Fig. 2).

```
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ make
make -C /usr/src/linux-5.5.8/ M=/home/galaxies/CS307-Projects/Project2/pid modules
make[1]: Entering directory '/usr/src/linux-5.5.8'
Building modules, stage 2.
MODPOST 1 modules
make[1]: Leaving directory '/usr/src/linux-5.5.8'
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo dmesg -C
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo insmod pid.ko
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo dmesg
[19487.271845] /proc/pid is created!
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ echo "1395" > /proc/pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ cat /proc/pid
command = [gmain] pid = [1395] state = [1]
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ echo "1" > /proc/pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ cat /proc/pid
command = [systemd] pid = [1] state = [1]
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ echo "5" > /proc/pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ cat /proc/pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo dmesg
[19487.271845] /proc/pid is created!
[19523.924274] Invalid PID!
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ echo "6" > /proc/pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ cat /proc/pid
command = [kworker/0:0H] pid = [6] state = [1026]
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo rmmod pid
galaxies@ubuntu:~/CS307-Projects/Project2/pid$ sudo dmesg
[19487.271845] /proc/pid is created!
[19523.924274] Invalid PID!
[19555.988406] /proc/pid is removed!
```

Figure 2: The execution result of my Linux Kernel module for task information

□

### 3 Personal Thoughts

During the first UNIX Shell project, I've experienced the process of implementing a simple UNIX Shell program. The implementation enhances my understandings of the pipe and input/output redirections. Actually the implementations of the UNIX Shell program is complicated than I thought before, which takes me about 4 hours to finish, and the amount of code reaches 8 KB. After finishing the project, I feel very fulfilled and I become more familiar with the Linux C instructions. I'm also getting well with the C language programming in Linux.

The Linux kernel module for task information project enhances my understandings of `/proc` file system, and it strengthen my knowledge about it, which is important in Linux system.

By the way, you can find all the source codes in the "src" folder. You can also refer to [my github](#) to see my codes of this project, and they are in the Project2 folder.