

# Project 8: Designing Virtual Memory Manager

CS307-Operating System (D), CS356-Operating System Course Design, Chentao Wu, Spring 2020.

Name: 方泓杰(Hongjie Fang) Student ID: 518030910150 Email: galaxies@sjtu.edu.cn

## 1 Designing Virtual Memory Manager

### 1.1 Requirements

This project consists of writing a program that translates logical to physical addresses for a virtual address space of size  $2^{16} = 65536$  bytes. The program will read from a file containing logical address and, using a TLB and a page table, will translate each logical address to its corresponding physical address and output the value of the byte stored at the translated physical address. It requires us to use simulation to understand the steps involved in translating logical address to physical address, and this will include resolving page faults using demand paging, managing a TLB, and implementing a page-replacement algorithm.

The program will read a file containing several 32-bit integer numbers that represent logical addresses. However, the 16-bit addresses is the only thing that needs to be concerned, so we must mask the rightmost 16 bits of each logical addresses. These 16 bits are divided into an 8-bit page number and an 8-bit page offset. Hence, the addresses are structured as shown as Fig. 1.




Figure 1: The address structure

Other specifics includes the following:

- $2^8$  entries in the page table;
- Page size of  $2^8$  bytes;
- 16 entries in the TLB;
- Frame size of  $2^8$  bytes;
- 256 frames;
- Physical memory of 65536 bytes ( $256 \text{ frames} \times 256\text{-byte frame size}$ ).

The program is to output the following values:

- The logical address being translated (the integer value being read from `addresses.txt`).
- The corresponding physical address (what your program translates the logical address to).
- The signed byte value stored in physical memory at the translated physical address.

We also provide the file `correct.txt`, which contains the correct output values for the file `addresses.txt`. You should use this file to determine if the program is correctly translating logical to physical addresses.

After completion, the program is to report the following statistics:

- Page-fault rate - The percentage of address references that resulted in page faults.
- TLB hit rate - The percentage of address references that were resolved in the TLB.

Since the logical addresses in `addresses.txt` were generated randomly and do not reflect any memory access locality, do not expect to have a high TLB hit rate.

After that, we need to change the frame number to 128 to implement the page replacement algorithm.

## 1.2 Methods

Here are some specific methods of the virtual memory manager.

- We implement a free frame list using the linked list, and at first all the frames are free. Every time we need a free frame, we can look it up in the free frame list using `get_empty_frame()` function. If the function returns `-1`, then there is no empty frame and the memory will need a page replacement; otherwise the return value of the function is the empty frame number.
- We implement the memory using the LRU page replacement strategy. Every time we need to add a page into the memory, we will call the `get_empty_frame()` function to get the free frame number. If there is a free frame, we will use it to store the page; otherwise we will use the LRU page replacement algorithm to find the victim page and replace it with the requested page, and we will use the `delete_page_table_item()` function to delete the corresponding page table item and TLB item.
- We implement the TLB also using the LRU replacement strategy. Every time we need to add an entry into the TLB, we will find out whether there is a free entry. If so, we can directly put it into the entry; otherwise we must replace one entry using the LRU algorithm. Every time TLB hits, we will also update the LRU value of each entry.
- We also implement the page table using the normal methods. Every time we need to transform the virtual address into the physical address, we will first look it up in TLB by calling `get_TLB_frame_num()` function provided by TLB. If TLB miss, we will refer to the page table. If page table miss, we will using the demand paging algorithm to put the page in the backing store into the memory. So the transforming process is as follows (Fig. 2).

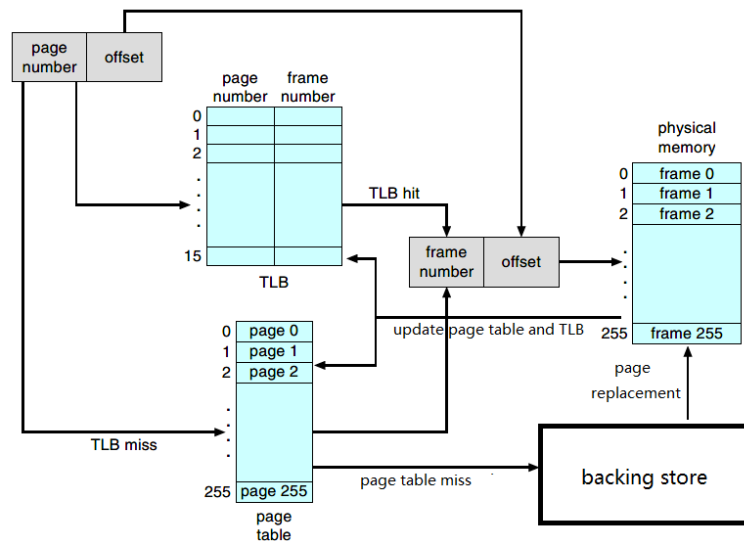


Figure 2: The transforming process

## 1.3 Implementation

Here is the specific implementation of the virtual memory manager (`vmm.c`).

[Note: we directly implement the page-replacement algorithm here.](#)

```
1 # include <stdio.h>
2 # include <stdlib.h>
3 # include <string.h>
4
5 # define PAGE_NUM 256
6 # define PAGE_SIZE 256
7 # define FRAME_NUM 128
8 # define FRAME_SIZE 256
9 # define TLB_SIZE 16
10
11 // ===== Empty Frame List ===== //
12 struct empty_frame_list_node {
13     int frame_num;
14     struct empty_frame_list_node *nxt;
15 };
16
17 struct empty_frame_list_node *head = NULL;
18 struct empty_frame_list_node *tail = NULL;
19
20 // Add the empty frame to the empty frame list.
21 void add_empty_frame(int frame_num) {
22     if (head == NULL && tail == NULL) {
23         tail = (struct empty_frame_list_node *) malloc (sizeof(struct
24             empty_frame_list_node));
25         tail -> frame_num = frame_num;
26         tail -> nxt = NULL;
27         head = tail;
28     } else {
29         tail -> nxt = (struct empty_frame_list_node *) malloc (sizeof(struct
30             empty_frame_list_node));
31         tail -> nxt -> frame_num = frame_num;
32         tail -> nxt -> nxt = NULL;
33         tail = tail -> nxt;
34     }
35 }
36
37 // Get an empty frame from the empty frame list.
38 // If success, return frame_num; otherwise, return -1.
39 int get_empty_frame() {
40     if (head == NULL && tail == NULL) return -1;
41
42     int frame_num;
43     if (head == tail) {
44         frame_num = head -> frame_num;
45         free(head);
46         head = tail = NULL;
```

```

45     return frame_num;
46 }
47
48 struct empty_frame_list_node *tmp;
49 frame_num = head -> frame_num;
50 tmp = head;
51 head = head -> nxt;
52 free(tmp);
53 return frame_num;
54 }
55
56 // Initialize the empty frame list.
57 void initialize_empty_frame_list() {
58     for (int i = 0; i < FRAME_NUM; ++ i)
59         add_empty_frame(i);
60 }
61
62 // Clean the empty frame list.
63 void clean_empty_frame_list() {
64     if (head == NULL && tail == NULL) return;
65     struct empty_frame_list_node *tmp;
66     while (head != tail) {
67         tmp = head;
68         head = head -> nxt;
69         free(tmp);
70     }
71     free(head);
72     head = tail = NULL;
73 }
74 // ===== End of Empty Frame List ===== //
75
76
77 // ===== Memory ===== //
78 char memory[FRAME_NUM * FRAME_SIZE];
79 int frame_LRU[FRAME_NUM];
80 char buf[FRAME_SIZE];
81 FILE *fp_backing_store;
82
83 void initialize_memory() {
84     fp_backing_store = fopen("BACKING_STORE.bin", "rb");
85     if (fp_backing_store == NULL) {
86         fprintf(stderr, "[Err] Open backing store file error!\n");
87         exit(1);
88     }
89     initialize_empty_frame_list();
90     for (int i = 0; i < FRAME_NUM; ++ i)
91         frame_LRU[i] = 0;
92 }
93
94 void delete_page_table_item(int frame_num);

```

```

95 int add_page_into_memory(int page_num) {
96     fseek(fp_backing_store, page_num * FRAME_SIZE, SEEK_SET);
97     fread(buf, sizeof(char), FRAME_SIZE, fp_backing_store);
98
99     int frame_num = get_empty_frame();
100     if (frame_num == -1) {
101         // LRU replacement.
102         for (int i = 0; i < FRAME_NUM; ++ i)
103             if (frame_LRU[i] == FRAME_NUM) {
104                 frame_num = i;
105                 break;
106             }
107         delete_page_table_item(frame_num);
108     }
109
110     for (int i = 0; i < FRAME_SIZE; ++ i)
111         memory[frame_num * FRAME_SIZE + i] = buf[i];
112     for (int i = 0; i < FRAME_NUM; ++ i)
113         if (frame_LRU[i] > 0) ++ frame_LRU[i];
114     frame_LRU[frame_num] = 1;
115     return frame_num;
116 }
117
118 char access_memory(int frame_num, int offset) {
119     char res = memory[frame_num * FRAME_SIZE + offset];
120     for (int i = 0; i < FRAME_NUM; ++ i)
121         if (frame_LRU[i] > 0 && frame_LRU[i] < frame_LRU[frame_num])
122             ++ frame_LRU[i];
123     frame_LRU[frame_num] = 1;
124     return res;
125 }
126
127 void clean_memory() {
128     clean_empty_frame_list();
129     fclose(fp_backing_store);
130 }
131 // ===== End of Memory ===== //
132
133
134 // ===== TLB ===== //
135 int TLB_page[TLB_SIZE], TLB_frame[TLB_SIZE];
136 int TLB_LRU[TLB_SIZE];
137 int TLB_hit_count;
138
139 void initialize_TLB() {
140     TLB_hit_count = 0;
141     for (int i = 0; i < TLB_SIZE; ++ i) {
142         TLB_page[i] = 0;
143         TLB_frame[i] = 0;
144         TLB_LRU[i] = 0;

```

```

145     }
146 }
147
148 // Get the corresponding frame number from TLB.
149 // Return non-negative number for the corresponding frame number;
150 // Return -1 for TLB miss.
151 // Note: it's needless to check the validation of page_num again.
152 int get_TLB_frame_num(int page_num) {
153     int pos = -1;
154     for (int i = 0; i < TLB_SIZE; ++ i)
155         if (TLB_LRU[i] > 0 && TLB_page[i] == page_num) {
156             pos = i;
157             break;
158         }
159
160     if (pos == -1) return -1;
161
162     // TLB hit.
163     ++ TLB_hit_count;
164     for (int i = 0; i < TLB_SIZE; ++ i)
165         if (TLB_LRU[i] > 0 && TLB_LRU[i] < TLB_LRU[pos])
166             ++ TLB_LRU[i];
167     TLB_LRU[pos] = 1;
168     return TLB_frame[pos];
169 }
170
171 // Update TLB entry
172 void update_TLB(int page_num, int frame_num) {
173     int pos = -1;
174     for (int i = 0; i < TLB_SIZE; ++ i)
175         if(TLB_LRU[i] == 0) {
176             pos = i;
177             break;
178         }
179
180     if (pos == -1) {
181         // LRU replacement.
182         for (int i = 0; i < TLB_SIZE; ++ i)
183             if(TLB_LRU[i] == TLB_SIZE) {
184                 pos = i;
185                 break;
186             }
187     }
188
189     TLB_page[pos] = page_num;
190     TLB_frame[pos] = frame_num;
191     for (int i = 0; i < TLB_SIZE; ++ i)
192         if (TLB_LRU[i] > 0) ++ TLB_LRU[i];
193     TLB_LRU[pos] = 1;
194 }

```

```

195
196 // Delete TLB item.
197 void delete_TLB_item(int page_num, int frame_num) {
198     int pos = -1;
199     for (int i = 0; i < TLB_SIZE; ++ i)
200         if(TLB_LRU[i] && TLB_page[i] == page_num && TLB_frame[i] == frame_num) {
201             pos = i;
202             break;
203         }
204
205     if (pos == -1) return;
206
207     for (int i = 0; i < TLB_SIZE; ++ i)
208         if (TLB_LRU[i] > TLB_LRU[pos]) -- TLB_LRU[i];
209     TLB_LRU[pos] = 0;
210 }
211 // ===== End of TLB ===== //
212
213
214 // ===== Page Table ===== //
215 int page_table[PAGE_NUM];
216 int vi_page_table[PAGE_NUM]; // vi: valid-invalid
217 int page_fault_count;
218
219 void initialize_page_table() {
220     page_fault_count = 0;
221     for (int i = 0; i < PAGE_NUM; ++ i) {
222         page_table[i] = 0;
223         vi_page_table[i] = 0;
224     }
225 }
226
227 // Get the corresponding frame number.
228 // Return non-negative number for the corresponding frame number;
229 // Return -1 for invalid page number.
230 int get_frame_num(int page_num) {
231     if (page_num < 0 || page_num >= PAGE_NUM) return -1;
232
233     int TLB_res = get_TLB_frame_num(page_num);
234     if (TLB_res != -1) return TLB_res;
235
236     if (vi_page_table[page_num] == 1) {
237         update_TLB(page_num, page_table[page_num]);
238         return page_table[page_num];
239     } else {
240         // Page fault.
241         ++ page_fault_count;
242         page_table[page_num] = add_page_into_memory(page_num);
243         vi_page_table[page_num] = 1;
244         update_TLB(page_num, page_table[page_num]);

```

```

245     return page_table[page_num];
246 }
247 }
248
249 // Delete page table item
250 void delete_page_table_item(int frame_num) {
251     int page_num = -1;
252     for (int i = 0; i < PAGE_NUM; ++ i)
253         if(vi_page_table[i] && page_table[i] == frame_num) {
254             page_num = i;
255             break;
256         }
257     if (page_num == -1) {
258         fprintf(stderr, "[Err] Unexpected Error!\n");
259         exit(1);
260     }
261     vi_page_table[page_num] = 0;
262     delete_TLB_item(page_num, frame_num);
263 }
264 // ===== End of Page Table ===== //
265
266
267 void initialize() {
268     initialize_page_table();
269     initialize_TLB();
270     initialize_memory();
271 }
272
273 void clean() {
274     clean_memory();
275 }
276
277
278 int main(int argc, char *argv[]) {
279     if (argc != 2) {
280         fprintf(stderr, "[Err] Invalid input!\n");
281         return 1;
282     }
283
284     FILE *fp_in = fopen(argv[1], "r");
285     if(fp_in == NULL) {
286         fprintf(stderr, "[Err] File Error!\n");
287         return 1;
288     }
289
290     FILE *fp_out = fopen("output.txt", "w");
291     if (fp_out == NULL) {
292         fprintf(stderr, "[Err] File Error!\n");
293         return 1;
294     }

```



```

295
296 initialize();
297
298 int addr, page_num, offset, frame_num, res, cnt = 0;
299 while(~fscanf(fp_in, "%d", &addr)) {
300     ++ cnt;
301     addr = addr & 0x0000ffff;
302     offset = addr & 0x000000ff;
303     page_num = (addr >> 8) & 0x000000ff;
304     frame_num = get_frame_num(page_num);
305     res = (int) access_memory(frame_num, offset);
306     fprintf(fp_out, "Virtual address: %d Physical address: %d Value: %d\n", addr,
307             (frame_num << 8) + offset, res);
308 }
309
310 fprintf(stdout, "[Statistics]\n TLB hit rate: %.4f %%\n Page fault rate: %.4f
311             %%\n", 100.0 * TLB_hit_count / cnt, 100.0 * page_fault_count / cnt);
312
313 clean();
314 fclose(fp_in);
315 fclose(fp_out);
316 return 0;
317 }

```

## 1.4 Testing

I write a Makefile file and a `checker.c` program to help testing the program. You can modify Line 7 in `vmm.c` to change the total frame number to test the page replacement algorithm, and the default number is 128. Since the physical address may be different because of the different replacement strategy, we only check the value stored in the virtual address. The `checker.c` will output **Accept** if the program output matches the correct output, and the checker will output **Wrong Answer** if the outputs do not match. Here is the specific implementation of the checker (`checker.c`).

```

1 # include <stdio.h>
2 # include <string.h>
3
4 int main() {
5     FILE *fp_out = fopen("output.txt", "r");
6     FILE *fp_ans = fopen("correct.txt", "r");
7
8     int accept = 1;
9     int right_val, out_val, cnt = 0;
10    while (~fscanf(fp_ans, "Virtual address: %d Physical address: %d Value: %d\n",
11                  &right_val)) {
12        ++ cnt;
13        if (fscanf(fp_out, "Virtual address: %d Physical address: %d Value: %d\n", &
14                  out_val) == EOF) {
15            printf("File length mismatch!\n");
16            accept = 0;
17            break;
18        }
19    }
20    if (accept) {
21        printf("Accept\n");
22    } else {
23        printf("Wrong Answer\n");
24    }
25    return 0;
26 }

```

```

16     }
17     if (right_val != out_val) {
18         printf("Error on line %d.\n", ++ cnt);
19         accept = 0;
20     }
21 }
22
23 if (accept == 0) printf("Wrong answer.\n");
24 else printf("Accept.\n");
25
26 fclose(fp_out);
27 fclose(fp_ans);
28 return 0;
29 }

```

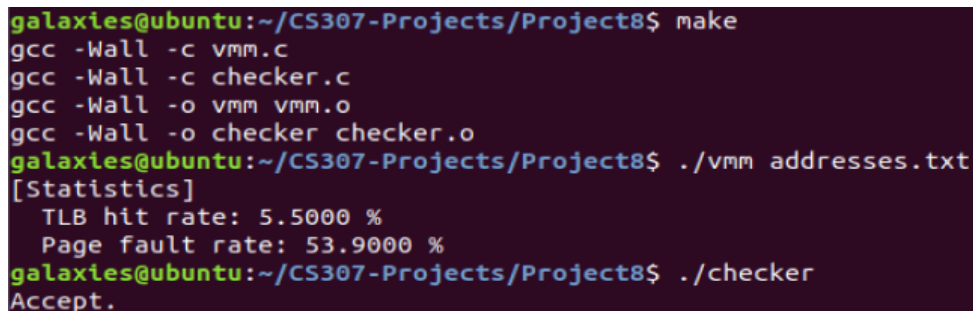
We only need to enter the following instructions in the terminal and we can begin testing.

```

1 make
2 ./vmm addresses.txt
3 ./checker

```

When the memory frame number is 128, the execution result of the virtual memory manager is as follows (Fig. 3).



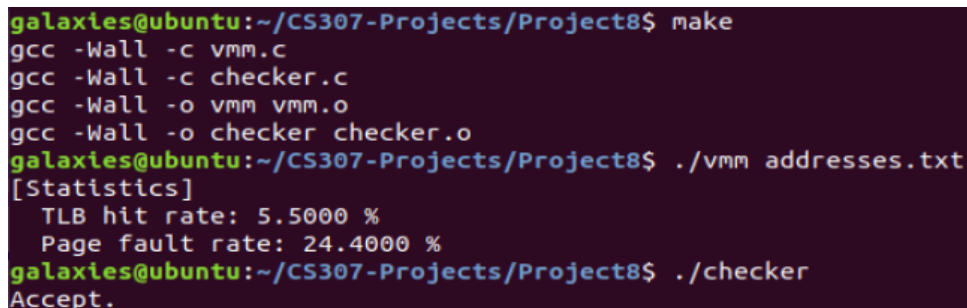
```

galaxies@ubuntu:~/CS307-Projects/Project8$ make
gcc -Wall -c vmm.c
gcc -Wall -c checker.c
gcc -Wall -o vmm vmm.o
gcc -Wall -o checker checker.o
galaxies@ubuntu:~/CS307-Projects/Project8$ ./vmm addresses.txt
[Statistics]
  TLB hit rate: 5.5000 %
  Page fault rate: 53.9000 %
galaxies@ubuntu:~/CS307-Projects/Project8$ ./checker
Accept.

```

Figure 3: The execution result of the virtual memory manager when memory frame number is 128

When the memory frame number is 256, the execution result of the virtual memory manager is as follows (Fig. 4).



```

galaxies@ubuntu:~/CS307-Projects/Project8$ make
gcc -Wall -c vmm.c
gcc -Wall -c checker.c
gcc -Wall -o vmm vmm.o
gcc -Wall -o checker checker.o
galaxies@ubuntu:~/CS307-Projects/Project8$ ./vmm addresses.txt
[Statistics]
  TLB hit rate: 5.5000 %
  Page fault rate: 24.4000 %
galaxies@ubuntu:~/CS307-Projects/Project8$ ./checker
Accept.

```

Figure 4: The execution result of the virtual memory manager when memory frame number is 256

When the memory frame number decreases, the page fault rate increases, which is correct since the memory can contain fewer frames.

## 2 Personal Thoughts

The project helps me understand the virtual memory manager better, including the mapping between virtual memory and physical memory, the replacement algorithm in TLB and memory and the backing store. The project also improves my understanding of page table and TLB - which is the cache of the page table. The project trains our coding skills of designing appropriate cache. The implementation of the algorithms is quite simple but it needs patience, and I enjoy the process of writing this program.

By the way, you can find all the source codes in the “src” folder. You can also refer to [my github](#) to see my codes of this project, and they are in the **Project8** folder.